

UCLA Computer Science 33 (Spring 2015)

Midterm

108 minutes total, open book, open notes

Questions are equally weighted (12 minutes each)

Name: _____ Student ID: _____

1	2	3	4	5	6	7	8	9	sum

1. Suppose you want a thread T1 to wait until thread T2 finishes, and that T2 is detached. Explain how to arrange for this reliably, assuming the threads cooperate by executing code that you specify. Your code may invoke any of the `pthread_*` or `sem_*` primitives discussed in the book or in class. Explain any assumptions you make and any race conditions that you couldn't fix (these should be reasonable and few).

2. Suppose we extend the bitwise operations \wedge , $\&$, $|$ and \sim to operate on floating-point values by applying these bitwise operations to their representations. For example, since the binary representation of `0.1f` is `0x3dcccccd`, and $\sim 0x3dcccccd == 0xc2333332$, and `0xc2333332` represents `-44.799995f`, then $\sim 0.1f$ would yield `-44.799995f` and $\sim -44.799995f$ would yield `0.1f`. Recall that the general rule for floating point operations is that NaNs are infectious, i.e., that if one or both inputs to a floating-point operation is a NaN, then the operation yields a NaN. Which (if any) of the bitwise operations \wedge , $\&$, $|$ and \sim are infectious on NaNs? Explain.

3. The function `stricmp(A, B)` compares the two strings A and B ignoring case, and returns an int. If we let `a` = the lowercased version of A and `b` = the lowercased version of B, then `stricmp(A, B)` returns a negative number if `a` compares less than `b`, a positive number if `a` compares greater than `b`, and zero otherwise. This function compares strings byte by byte, and assumes only the 52 ASCII letters.

Consider the following `stricmp` implementation. Assume that it's running on the x86.

```
#include <string.h>

#define min(a, b) ((a) < (b) ? (a) : (b))

char
cvtlower (char c)
{
    if ('A' <= c && c <= 'Z')
        return c - 'A' + 'a';
    return c;
}
```

```

}

int
strcmp (char const *a, char const *b)
{
    for (size_t i = 0;
         i < min(strlen(a), strlen(b));
         i++)
        if (cvtlower(a[i]) < cvtlower(b[i]))
            return -1;
        else if (cvtlower(a[i]) > cvtlower(b[i]))
            return 1;
    return 0;
}

```

Propose two optimizations of this code, at least one of which is likely to improve performance greatly and the other at least somewhat. Explain why the former is likely to be better than the latter.

4. Consider the following C function and its translation to x86-64 assembly language. The C function returns `a[i]` in the typical case where `i` is in range, and returns 0 otherwise:

```

int
subscript (int *a, unsigned i,
           unsigned int n)
{
    if (0 <= i && i < n)
        return a[i];
    return 0;
}

subscript:
    xorl    %eax, %eax
    cmpl    %edx, %esi
    jnb     .L2
    movl    %esi, %esi
    movl    (%rdi,%rsi,4), %eax
.L2:
    ret

```

4a. How can this code be correct? The source has two comparisons, but the assembler has just one.

4b. Why aren't conditional moves helpful for improving this code's performance? Explain.

5. Suppose your program has three parts that are done in sequence and take 0.5, 0.3, and 0.2 of the time respectively. You can parallelize the first part and speed it up by a factor of 2. Or you can parallelize the second part and speed it up by a factor of 8. Use Amdahl's law to calculate which of these two will give you better performance and why. Suppose you can do both

parallelizations: how much will your performance improve compared to the original, or to either parallelization alone? Show your work.

6. Let d = the number 0.1 in C (i.e., the 'double' value 0.1). Let f = the number 0.1f in C (i.e., the 'float' value 0.1f). And let r = the number 0.1 in mathematics (i.e., the real number equal to 1 divided by 10). Recall that the binary representation of r is the repeating sequence 0.000110011001100110011... base 2. Sort the values d , f , and r into nondescending order. If two or more of these three values are equal, say so. Assume x86-64 arithmetic with default rounding. Show your work.

7. We have a special kind of SRAM cache called Cache Z. Cache Z uses a write-back approach, but omits the dirty bit. Instead, whenever it needs to know whether a cache line is dirty, it loads the corresponding data from RAM and compares it to the data in the cache line: if they compare equal, Cache Z acts as if the dirty bit were zero, and if they compare nonequal, Cache Z acts as if the dirty bit were nonzero. Compare the pros and cons of Cache Z to an ordinary write-back write-allocate cache. Propose an improvement to Cache Z's performance that does not involve adding a dirty bit.

8. Suppose you have an x86-64 machine with the following characteristics:

```
two sockets
12 CPU cores per socket
L1 instruction cache: 32 KiB per core
L1 data cache: 32 KiB per core
L2 cache: 256 KiB per core
L3 cache: 30 MiB per socket
64 GiB DRAM per socket
```

L1 and L2 caches are private to each core.
L3 cache is shared by all cores in a socket.
All caches are writeback.

8a. The single instruction cache at L1 is fast but very small: why won't performance suffer greatly if your program's kernel doesn't fit into 32 KiB?

8b. Consider the following program, and assume its x86-64 code fits entirely within the L1 instruction cache, and assume that the source and destination do not overlap.

```
#define N <<you pick the constant>>
void transpose(int dst[N][N], int src[N][N]) {
    for (int i = 0; i < N; i++)
        for (int j = 0; j < N; j++)
            dst[j][i] = src[i][j];
}
```

Suppose this function is often executed in your multithreaded application on the specified machine. What values of N do you recommend for

good performance, and why? Look for local sweet spots for N . State any further assumptions you're making.

9. Section 5.9.2 of the book shows how to apply the associative law to improve performance while unrolling a loop. Can we use a similar idea to improve performance by applying the distributive law $A*(B + C) == A*B + A*C$? Why or why not?